

GSoC 2026 Proposal (R Project): Time-Dependent Constraints in gfpop

Weidong “William” Zhang

March 31, 2026

Contents

1	Logistics	2
1.1	Project Info	2
1.2	Contact Information	2
1.3	Contributor Affiliation	2
1.4	Schedule Conflicts	2
1.5	Mentors	2
1.6	Qualification Tests	4
2	Project Plan	6
2.1	Overview	6
2.2	gfpop Architecture	6
2.3	The Problem: Why gfpop Cannot Express Label Constraints	7
2.4	The Extension: A Rule Function	9
2.5	What the User Sees: The R API	10
2.6	What the Solver Does: C++ Implementation	12
2.7	Obstacles and Challenges	14
2.8	Deliverables and Validation	15
2.9	Success Criteria and Non-Goals	17
2.10	Timeline	18
2.11	Management of Coding Project	19
3	About Me	20
3.1	Bio	20
3.2	Programming Languages and Tools	20
3.3	Skills I Picked Up	20
3.4	What I Expect from GSoC	21
4	References	22

1 Logistics

1.1 Project Info

- **Project title:** Time-Dependent Constraints in gfpop
- **Project short title:** Time-dep constraints in gfpop
- **Project size:** Large (350 hours)
- **URL of project idea page:**
<https://github.com/rstats-gsoc/gsoc2026/wiki/time-dependent-constraints-in-gfpop>

1.2 Contact Information

- **Contributor name:** Weidong “William” Zhang
- **Postal address:** T3L 1Z3
- **Telephone:** (825) 365-2426
- **Email:** weidongzhang60@gmail.com
- **GitHub:** [@williamzhang7792](https://github.com/williamzhang7792)
- **LinkedIn:** [weidong-william-zhang](https://www.linkedin.com/in/weidong-william-zhang)
- **Discord username:** weidong.z

1.3 Contributor Affiliation

- **Institution:** University of Waterloo
- **Program:** Bachelor of Mathematics, Statistics major, Computing minor
- **Stage of completion:** Second year
- **Contact to verify:** mathuo@uwaterloo.ca

1.4 Schedule Conflicts

I will be taking 3–5 summer courses at my university, depending on offerings. I have no other jobs, internships, or commitments. I plan to dedicate approximately 25–35 hours per week to GSoC alongside my coursework, with heavier weeks (35+) in June before course deadlines peak and lighter weeks (25) in August. To compensate for the split workload, I plan to front-load work during the community bonding period (R-side API, design doc, dev environment) so that coding starts with a running head start.

1.5 Mentors

- **Evaluating mentor:**
 - Vincent Runge, vincent.runge@univ-evry.fr
- **Co-mentors:**
 - Toby Dylan Hocking, tdhock5@gmail.com
 - Gaetano Romano, g.romano@lancaster.ac.uk

Have you been in touch with the mentors? When and how?

Yes. I have been in contact with all three mentors via email and GitHub since late February 2026.

- **Feb ~24:** Emailed Vincent Runge (CC'd Toby Hocking and Gaetano Romano) introducing myself and sharing completed Easy, Medium, and Hard test results. Asked for feedback on my understanding of the FPOP algorithm.
- **Feb ~26:** Toby reviewed the test results and gave feedback on code readability (unicode escapes in R). He pointed out that the solvers should be integrated into a proper PeakSegOptimal fork rather than standalone scripts, and suggested I begin the GSoC application. I refactored both solvers into the fork and updated Toby on March 2.
- **Mar ~10:** Toby suggested I submit a PR so the diff could be reviewed. Opened [PR #24](#) (add unconstrained Poisson FPOP and isotonic Normal FPOP solvers) on PeakSegOptimal.
- **Mar 11–27:** Opened [PR #25](#) to migrate CI from Travis to GitHub Actions. Over the course of several iterations, Toby helped me work through CI configuration, R packaging practices (`Suggests`, `Remotes`, `extra-packages`), test dependency management for archived CRAN packages, and a cache invalidation issue caused by an Ubuntu runner update. PR **merged** on Mar 27.
- **Mar ~15:** Emailed Vincent directly with the draft proposal. Vincent responded with technical feedback: the $O(n \log n)$ complexity claim for FPOP/GFPOP is not proven (he prefers “quasi-linear”), and backtracking under constraints is a known difficulty. I incorporated both points into the proposal.

The review process has been a good learning experience. Through [PR #25](#), I picked up practical R packaging knowledge (how `setup-r-dependencies` resolves `Suggests`, how `Remotes` interacts with CRAN checks) that I would not have learned from the documentation alone. Also, thanks to Vincent's feedback, I was able to improve the technical claims in the proposal.

1.6 Qualification Tests

All three tests completed, as described on the [project wiki](#). Code is in the [test repository](#); both solvers are also integrated into a [PeakSegOptimal fork](#).

1.6.1 plotModel(graph) Visualization (Easy Test)

I wrote an R function that takes a gfpop graph with a `rule` column and draws it as a state-by-rule transition matrix using ggplot2. Rows are states, columns are rules, arrows show allowed transitions. Self-loops are stacked vertically and edges are shortened to avoid label overlap.

This visualization already models the API extension. The `plotModel()` output in the R API section was generated by this function.





Figure 1: `plotModel()` output from the Easy test: LOPART graph with 4 rules and 2 states.

1.6.2 Unconstrained FPOP (Medium Test)

I adapted PeakSegOptimal’s 2-state constrained solver into a 1-state unconstrained solver for optimal partitioning with Poisson loss. The key insight was replacing `set_to_min_less_of` / `set_to_min_more_of` with a new `set_to_unconstrained_min_of` that collapses the cost to a single flat piece at the global minimum.

- 9 testthat tests, oracle-validated against `Segmentor3IsBack::Segmentor(model=1)` across multiple penalties (replacing with `gfpop::gfpop()` per mentor feedback in [PR #25](#)).
- Integrated into the PeakSegOptimal fork with C registration, `.C()` interface, and 28 tests.
- Prepared a [slide deck](#) walking through FPOP pruning mechanics.

$t = 2$: The Jump ($y_2 = 10$)



9 / 14

Figure 2: *FPOP pruning trace from the Medium test, showing the number of piecewise cost function intervals over time.*

1.6.3 Regularized Isotonic Regression (Hard Test)

I implemented `NormalLossPiece` with closed-form quadratic roots (no Newton iteration needed, unlike the Poisson case). Both `NormalLossPiece` and `PoissonLossPieceLog` inherit from a shared `LossPiece` base class, which is directly relevant to `gfpop`'s multi-loss architecture.

I also reconstructed `set_to_min_less_of` from first principles in a standalone setting. This is the core operator behind the isotonic constraint, and the same one that will enforce up-constraints in the time-dependent extension.

- 14+ testthat tests including a 200-trial fuzz test against `isoreg()` (PAVA).
- `penalty=0` matches `isoreg`; positive penalty matches `fpop::Fpop()` on non-decreasing data.




Figure 3: *Visualization of the `set_to_min_less_of` operator from the Hard test. The blue curve is the original piecewise quadratic cost; the red curve is the running minimum.*

2 Project Plan

2.1 Overview

The `gfpop` package (Runge et al., 2020) unifies a wide family of constrained changepoint problems (isotonic regression, up-down peak detection, and more) through a single graph abstraction. Combined with the FPOP pruning strategy (Maidstone et al., 2017), this yields an empirically quasi-linear solver that is both general and fast.

But the graph is *time-independent*: the same edges apply at every data point. This means `gfpop` cannot express models like LOPART (Hocking and Srivastava, 2022), where the allowed transitions depend on whether the current observation falls inside a labeled region. Today, R packages for changepoint detection split into two categories: domain-specific solvers (LOPART, FLOPART) that handle labels but only for one model, and generic frameworks (`gfpop`) that handle many models but not labels.

The extension to close this gap is small and surgical. A rule function $r : \{1, \dots, N\} \rightarrow \{1, \dots, R\}$ selects which subset of edges is active at each time step. The R API changes are minimal (one new parameter on `Edge()`, one new argument to `gfpop()`). The C++ changes are concentrated in a single loop but require careful infinity handling and edge-sort invariant preservation. Once in place, `gfpop` can express both LOPART and models that no current package can handle, like up-down peak detection with label constraints. For the R ecosystem, this means any future labeled changepoint model can be expressed in `gfpop` without writing a new package from scratch.

I have completed all three qualification tests (Easy, Medium, and Hard), as recommended by the project wiki: an unconstrained FPOP solver (Poisson loss), a regularized isotonic regression solver (Normal loss), and a rule-aware graph visualizer. That is roughly 1,200 lines of C++ exercising every component this extension will touch: the DP loop, the constraint operators, the min-envelope, and the backtracking.

I'm interested in this project because it involves some serious algorithm design and real scientific applications. Also, being able to ship something inside a mature CRAN package is quite cool.

2.2 `gfpop` Architecture

Before diving into the limitation and the fix, it helps to see how `gfpop` is structured.




Figure 4: *gfpop* pipeline from R to C++. The highlighted box is the single modification target for time-dependent constraints.

The pipeline:

1. The user constructs a graph via `graph()`, `Edge()`, `Node()`, `StartEnd()`.
2. `graphReorder()` sorts edges by (`state2`, `penalty`) and converts state names to 0-based integers.
3. `gfpopTransfer()` (the Rcpp bridge) reads the graph into C++ `Edge/Graph` objects and sets up the cost function dispatch via global function pointers.
4. `Omega::gfpop()` runs the DP with three steps per data point: `LP_edges_operators(t)` applies constraint operators to each edge’s source cost function; `LP_edges_addPointAndPenalty(t)` adds the data point cost and edge penalty; `LP_t_new_multipleMinimization(t)` takes the pointwise minimum across all incoming edges for each destination state.

The edge sort invariant (edges sorted by (`state2`, `penalty`)) is exploited by step 4’s minimization: it walks the edge array in order, processing all edges targeting the same state in one contiguous block. If edges with the same target state are not adjacent, this walk breaks. Any time-dependent filtering must preserve this invariant.

2.3 The Problem: Why gfpop Cannot Express Label Constraints

2.3.1 A Concrete Scenario

Consider a genomics researcher analyzing ChIP-seq data. They have expert-labeled regions indicating where peaks should (positive labels) and should not (negative labels) appear. They want a changepoint model that respects these labels, a model where the set of allowed transitions depends on the current position in the data. The LOPART package can do this, but only for Gaussian loss with no up-down constraints. *gfpop* can handle arbitrary losses and graph constraints, but not labels.

The [project wiki page](#) frames this as the central gap: *gfpop*’s graph is static, but real-world models need edges that switch on and off depending on context.

2.3.2 The Fragmented Landscape

The table below (adapted from Kaufman et al., 2024) maps the current algorithm landscape. No single generic framework handles both label constraints and up-down constraints:

Algorithm	Labels	Up-down	Time
OPART (Jackson et al., 2005)	No	No	$O(n^2)$
LOPART (Srivastava & Hocking, 2023)	Yes	No	$O(n^2)$
FPOP (Maidstone et al., 2017)	No	No	quasi-linear*
GFPOP (Runge et al., 2020)	No	Yes	quasi-linear*
FLOPART (Kaufman et al., 2024)	Yes	Yes	quasi-linear*

*Empirically quasi-linear; worst-case $O(n^2)$ since pruning is not guaranteed.

Table 1: *Algorithm landscape for changepoint detection (adapted from Kaufman et al., 2024, Table 1). GFPOP is the most general framework but cannot use labels; LOPART uses labels but has no up-down constraints and runs in $O(n^2)$. FLOPART bridges the gap, but only for one model.*

As Kaufman et al. (2024) write: “The current fastest constrained algorithm, GFPOP, can be considered unsupervised because the algorithm does not use the labels in the train set, allowing for train errors. In contrast, changepoint models that currently utilize labels are all unconstrained models.” The FLOPART CRAN vignette demonstrates this empirically: on labeled ChIP-seq data, every unlabeled model, regardless of penalty, produces either false positives or false negatives. Only the labeled model achieves zero label errors.

Why not just use FLOPART? FLOPART solves one specific model: Poisson loss with up-down constraints and labels. It is a standalone C++ solver, not part of gfpop’s graph framework. Any new combination (Gaussian loss with isotonic constraints and labels, Huber loss with up-down constraints and labels) would require writing yet another standalone solver from scratch. The [project wiki page](#) frames this clearly: “Existing R packages fall into two categories: domain-specific packages such as LOPART ... generic frameworks such as gfpop ... (but do not allow time dependent constraints).” Adding time-dependent constraints to gfpop means every future model gets labels for free. One extension instead of N standalone solvers.

2.3.3 The Current gfpop Objective

The gfpop framework minimizes (see Runge et al., 2020 for full derivation):

$$Q_n = \min_{\substack{p \in \mathcal{G}_n \\ \mu | p(\mu)}} \sum_{t=1}^n [\gamma_{e_t}(y_t, \mu_t) + \beta_{e_t}]$$

In words: find the path through the constraint graph and the segment means that minimize total loss plus total penalty. Here $\mathcal{G}_n$ is the set of valid paths through a *collapsed* constraint graph G , γ_{e_t} is the loss function on edge e_t , $\beta_{e_t} \geq 0$ is the penalty, and the indicator constraints $I_{e_t}(\mu_t, \mu_{t+1}) = 1$ enforce relationships between consecutive segment means (e.g., non-decreasing for “up” edges).

The key limitation: the same graph G applies at every time step t .

2.3.4 The LOPART Objective

LOPART extends optimal partitioning by restricting the set of candidate last-changepoints T_t at each time step based on label positions (see Srivastava and Hocking, 2023 for the full formulation). The update rule has four cases depending on where data point t falls:

- **Unlabeled:** $t-1$ is added as a changepoint candidate. The solver is free to place a change here or not.
- **Inside a positive label:** no new candidate is added. The solver cannot introduce a *new* changepoint, but candidates from earlier time steps remain available.
- **End of a positive label:** the candidate set is reset to positions *inside* the label, forcing at least one changepoint within it.
- **Inside a negative label:** no new candidate is added and the solver is blocked from placing any changepoint here.

Each case allows a different set of transitions. This is precisely a time-dependent constraint, and gfpop cannot express it.

2.3.5 The Gap in Code

In `Omega::LP_edges_operators(t)` ([src/Omega.cpp](#)), gfpop iterates over *all* edges at *every* time step:

```
for (unsigned int i = 0; i < q; i++) {
    LP_edges[i].LP_edges_constraint(
        LP_ts[t][m_graph.getEdge(i).getState1()],
        m_graph.getEdge(i), t);
}
```

This is the single point in the codebase that makes gfpop time-independent (see Figure~4). The entire project reduces to making this loop conditional on the current time step.

2.4 The Extension: A Rule Function

The extension is minimal: instead of applying all edges at every time step, apply only the subset that matches the current data point’s rule.

Definition. Define a rule function $r : \{1, \dots, n\} \rightarrow \{1, \dots, R\}$ that maps each data point to a rule ID. Partition the edge set into R subsets $E_1, \dots, E_R$, where each edge belongs to the subset corresponding to its rule ID. Edges with rule = NA belong to *all* subsets (backward compatibility).

The modified DP. Let $O_t^{s,s'}(\theta)$ denote the result of applying the constraint operator of edge (s, s') (up, down, or identity) to the cost function Q_t^s . The update equation becomes (see Runge et al., 2020 for the standard gfpop formulation):

$$Q_{t+1}^{s'}(\theta) = \min_{s:(s,s') \in E_{r(t+1)}} \left\{ O_t^{s,s'}(\theta) + \gamma_{(s,s')}(y_{t+1}, \theta) + \beta_{s,s'} \right\}$$

The only change from standard gfpop: the minimization is over $E_{r(t+1)}$ instead of E .

The infinity requirement. When state s' has no incoming active edges under rule $r(t)$, the cost function must remain $Q_t^{s'}(\theta) = +\infty$ for all θ . Ensuring that the infinity sentinel propagates correctly through gfpop's linked-list cost functions is the hardest sub-problem of this project. Section~2.6 details exactly where infinity appears and how each function handles it.

Per-step edge processing is $O(q_{\text{active}})$ instead of $O(q)$, where q is the total number of edges and $q_{\text{active}} \leq q$ is the number active under the current rule. The FPOP pruning behavior is unchanged, so empirical complexity remains quasi-linear in N .

2.5 What the User Sees: The R API

2.5.1 Edge() Gains a rule Parameter

```
gfpop::Edge("normal", "normal", type = "null", rule = 1)
gfpop::Edge("normal", "normal", type = "std", rule = 1, penalty = 5.5)
```

Default: `rule = NA`, meaning the edge is active in all rules (backward compatible with all existing gfpop code).

2.5.2 gfpop() Gains a rule Argument

```
gfpop::gfpop(data.vec, mygraph, type = "mean", rule = rule_vec)
```

The `rule_vec` is an integer vector of length N , constructed from the label annotations:

```
LOPART.graph <- gfpop::graph(
  gfpop::Edge("normal", "normal", type = "null", rule = 1),
  gfpop::Edge("normal", "normal", type = "std", rule = 1, penalty = 5.5),
  gfpop::Edge("normal", "normal", type = "null", rule = 2),
  gfpop::Edge("noChange", "noChange", type = "null", rule = 2),
  gfpop::Edge("normal", "noChange", type = "std", rule = 2, penalty = 5.5),
  gfpop::Edge("normal", "normal", type = "std", rule = 3, penalty = 5.5),
  gfpop::Edge("noChange", "normal", type = "null", rule = 3),
  gfpop::Edge("normal", "normal", type = "null", rule = 4),
  gfpop::StartEnd(start = "normal", end = "normal"))

# user-side code; gfpop itself does not depend on data.table
rule_vec <- data.table::fcase(
  unlabeled, 1L,
  in.positive.label, 2L,
  end.positive.label, 3L,
  in.negative.label, 4L)

result <- gfpop::gfpop(data.vec, LOPART.graph, type = "mean", rule = rule_vec)
```

2.5.3 A Researcher's Workflow

The graph above encodes LOPART as 8 edges across 4 rules. The user builds a rule vector from their labels, calls `gfpop(..., rule = rule_vec)`, and gets labeled-optimal segments. No separate package, no format conversion, no loss-function restriction. Figure~5 shows how the active subgraph changes as the data moves through labeled regions.




Figure 5: The LOPART graph under each rule. Blue arrows are null (stay) edges; red arrows are change edges. Dashed gray edges are inactive under that rule. As the data moves through labeled regions, different edges switch on and off.

2.5.4 Notes on Compatibility and Validation

Backward compatibility. When `rule` is omitted, all edges are active at all times: current behavior is preserved. The `rule` column in the graph defaults to NA; NA-rule edges are always active. Existing user code is unaffected.

Input validation. `length(rule)` must equal `length(data)`. Rule IDs in `rule` must be a subset of rule IDs in the graph. `graphReorder()` must preserve the `rule` column and maintain the sort invariant within each rule group. The `weights` argument is independent of the rule vector (weights affect the loss, rules affect which edges are active), so no interaction is expected. Label boundary transitions (e.g., a positive label ending at t followed by a negative label at $t+1$) are handled by the rule vector, which assigns exactly one rule per data point. R-side validation will check that rule assignments cover all label boundaries without gaps or overlaps.

plotModel() already visualizes this. The `plotModel()` function from my Easy test already visualizes a rule-partitioned graph as a state-by-rule transition matrix:




Figure 6: *plotModel()* output for the LOPART graph above (same as Figure 1).

2.6 What the Solver Does: C++ Implementation

The C++ changes touch five files. The core modification is a conditional in `Omega.cpp`; the other changes support it.

2.6.1 Edge.h: Add ruleID

Add an `unsigned int ruleID` field and `getRuleID()` getter to the `Edge` class. The constructor reads it from the R data frame’s `rule` column. Edges with `rule = NA` receive a sentinel value (`UINT_MAX`) meaning “always active.”

2.6.2 Graph.h: Precomputed Rule-to-Edge Lookup

Two design options:

- **Option A (simple):** At each time step, loop over all q edges and check `edge.getRuleID() == rule_vec[t]`. $O(q)$ per step.
- **Option B (precomputed):** At construction, build a `std::vector<std::vector<unsigned int>>` mapping each rule ID to its sorted edge indices. At each time step, iterate only over `rule_to_edges[rule_vec[t]]`. $O(\text{active edges})$ per step.

I plan to start with Option A to validate the algorithm immediately, then profile. For the LOPART model ($q = 8$), an 8-iteration loop with a branch is likely faster than indexing a `std::vector<std::vector<unsigned int>>`. If profiling shows per-step overhead at larger q , Option B can be added as a targeted optimization. Either way, the lookup must preserve the (state2, penalty) sort order within each rule group so that `LP_t_new_multipleMinimization` continues to work correctly (it walks the edge array by target state; non-contiguous groups would break this walk).

2.6.3 Omega.cpp: LP_edges_operators(t), the Key Change

The modified function has three phases: (1)~reset all edge cost functions, (2)~apply constraint operators to active edges only, (3)~set remaining edges to the infinity sentinel.

```
// BEFORE (current gfpop)
for (unsigned int i = 0; i < q; i++) {
    LP_edges[i].LP_edges_constraint(
        LP_ts[t][m_graph.getEdge(i).getState1()],
        m_graph.getEdge(i), t);
}

// AFTER (with time-dependent rules)
for (unsigned int i = 0; i < q; i++)
    LP_edges[i].reset(); // deallocate old list
const auto& active = m_graph.activeEdges(rule_vec[t]);
for (unsigned int idx : active) {
    Edge const& edge = m_graph.getEdge(idx);
    LP_edges[idx].LP_edges_constraint( // reset() inside is a no-op
        LP_ts[t][edge.getState1()], edge, t);
}
for (unsigned int i = 0; i < q; i++) { // set remaining to infinity
    if (LP_edges[i].head == NULL)
        LP_edges[i].addFirstPiece(
            new Piece(Track(), interval, infinityCost));
}
```

Reset all edges first, then process only the active ones. Since `LP_edges_constraint` calls `reset()` internally, the second reset is a no-op (head is already NULL). The final loop sets unprocessed edges to the infinity sentinel. This avoids the double-allocation that would occur if `setInfinite()` allocated a piece only for `LP_edges_constraint` to immediately deallocate it. For larger graphs an alternative is to carry an `isActive` flag and skip inactive edges in all three DP sub-steps, avoiding allocation entirely. I plan to profile both approaches.

2.6.4 Omega.cpp: LP_t_new_multipleMinimization(t)

No structural change needed. The current code walks edges sorted by `state2` and takes the pointwise minimum via `LP_ts_Minimization`. Since `LP_ts[t+1][j]` is initialized to $+\infty$ and inactive `LP_edges` are also infinite, the `isInfinite()` short-circuit in `LP_ts_Minimization` will skip them. States with no active incoming edges naturally retain their infinite cost.

If all incoming edges for a state are inactive, the precomputed rule-to-edge lookup can detect this and skip the state entirely, avoiding repeated infinity-vs-infinity comparisons.

2.6.5 ListPiece.cpp: Infinity Handling in LP_ts_Minimization

Inactive edges are skipped in `LP_edges_operators`. In the proposed models, the rule design ensures that `operatorUp` and `operatorDw` are only applied to edges with reachable source states. As a safety

measure, I plan to add an `isInfinite()` guard in these operators as well (FLOPART's C++ source uses a similar short-circuit before each operator call). The simpler operators (null = copy, std = global-min) can see infinite source states and handle them trivially ($\text{copy}(\infty) = \infty$; $\text{min}(\infty) = \infty$). The infinity sentinel therefore only requires careful handling in two places:

1. `LP_ts_Minimization`: merges an infinite `LP_edges[k]` with a finite `LP_ts[t+1][j]` (or vice versa). Must return the finite side unchanged.
2. `LP_edges_addPointAndPenalty`: adds a data point cost to an infinite `LP_edges[i]`. For standard losses ($K = \infty$), the existing code adds finite values to the infinite cost constant, which stays infinite. For robust losses ($K \neq \infty$), the code performs interval root-finding that could produce NaN on an infinite cost. An `isInfinite()` short-circuit is needed for this path.

Add an `isInfinite()` predicate on `ListPiece` and short-circuit at the top of each function: if the incoming list is infinite, return immediately without modifying the target. This avoids pointer arithmetic on the sentinel piece and guards against root-finding on infinite costs.

An alternative design is to skip inactive edges in `LP_edges_addPointAndPenalty` as well (using the same active set), which eliminates the need for infinity handling there entirely. I plan to prototype both approaches and benchmark them.

2.6.6 main.cpp: `gfpopTransfer()` Bridge

Read the `rule` column from the graph `DataFrame` (as `IntegerVector`) and the `rule` argument from the function call. Pass both to the `Omega` constructor. If `rule` is not provided (R's default `NULL`), construct a vector of zeros and mark all edges as belonging to rule 0.

2.7 Obstacles and Challenges

Infinity handling in the min-envelope. The infinity sentinel analysis in Section~2.6 identifies `LP_ts_Minimization` and `LP_edges_addPointAndPenalty` as the two functions that need an `isInfinite()` short-circuit. I plan to write isolated test cases for each with infinite inputs before integrating into the full DP. Of the four obstacles below, infinity handling carries the most risk (novel code in a complex linked-list library); the others follow established modification patterns.

Edge sort invariant. The global edge order must remain `(state2, penalty)` because `LP_t_new_multipleMinimization` walks edges sequentially by `state2`. Changing the sort to `(rule, state2, penalty)` would split same-`state2` edges across rule groups, breaking this walk. Instead, `graphReorder()` keeps its current sort. The rule-to-edge lookup stores indices into the global array, grouped by rule. Each rule group's indices are already in `(state2, penalty)` order because they are a subset of the globally sorted array.

Backtracking through rule transitions. Backtracking under constraints is a known difficulty in constrained FPOP (flagged by Vincent Runge as well). The current backtracking code in `Omega::backtracking()` calls `m_graph.findBeta(state1, state2)` to recover the edge penalty at each step, looking up a unique edge by state pair. With rules, the same state pair can carry different penalties in different rules (e.g., `normal→normal` has $\beta = 5.5$ in Rule~1 but $\beta = 0$ in Rule~3). Two options: (a)~extend `findBeta` to accept a rule ID, recoverable via `rule_vec[t]`

since the **Track** already stores the time step $\sim t$; or (b)~store the edge penalty directly in **Track**, avoiding the lookup entirely. I plan to study how FLOPART handles its backtracking step before committing to an approach, and will discuss the trade-offs with the mentors early in the community bonding period.

Contingency. If the 6-rule up-down model stalls, the fallback is clear: ship the 4-rule LOPART model with full test coverage and benchmarks, and document the 6-rule design for a follow-up PR.

2.8 Deliverables and Validation

2.8.1 LOPART in gfpop (4 Rule IDs, 2 States)

States: **normal**, **noChange**. Each rule corresponds to a label context:

Rule	Label Context	Active Edges
1	Unlabeled	normal $\rightarrow$ normal (null), normal $\rightarrow$ normal (std + penalty)
2	In positive label	normal $\rightarrow$ normal (null), noChange $\rightarrow$ noChange (null), normal $\rightarrow$ noChange (std + penalty)
3	End of positive label	normal $\rightarrow$ normal (std + penalty), noChange $\rightarrow$ normal (null)
4	In negative label	normal $\rightarrow$ normal (null) only

Table 2: LOPART rule structure: 2 states $\times$ 4 rules. See Figure 6 for the `plotModel()` visualization.

Why this encoding reproduces LOPART. The **noChange** state tracks whether a changepoint has occurred inside the current positive label. During a positive label (Rule 2), the only way to reach **noChange** is via a changepoint from **normal**. At the label boundary (Rule 3), **noChange** $\rightarrow$ **normal** (null) carries the post-change cost forward, while **normal** $\rightarrow$ **normal** (std) forces a change for paths that haven't changed yet. Either way, at least one changepoint occurs inside every positive label, matching LOPART's T_t reset. In negative labels (Rule 4), only the null self-loop is active, so no changepoint can occur. The formal equivalence will be verified by oracle testing.

Validation: Oracle test against `LOPART::LOPART()`. On shared test data (Gaussian loss, same penalty), the gfpop output with rule vector must exactly match LOPART's output: same segments, same means, same cost.

2.8.2 Up-Down with Labels (6 Rule IDs, 4 States)

States: **Up**, **Dw**, **noChangeUp** (nCU), **noChangeDw** (nCD). The graph requires `StartEnd(start = c("Up", "Dw"), end = c("Up", "Dw"))` so that nCU and nCD begin at infinite cost (unreachable until a change occurs inside a positive label). This model does not exist in any current R package. gfpop would be the first to support it.

The 6 rules arise from splitting label contexts to avoid operating on infinite-cost functions at unreachable states. This is a deliberate design choice: instead of relying on correct infinity arithmetic (the hardest obstacle), the rule function prevents infinity from entering the operators in the first place.

Rule	Label Context	Active Edges
1	Unlabeled	Up $\rightarrow$ Up (null), Dw $\rightarrow$ Dw (null), Dw $\rightarrow$ Up (up + pen), Up $\rightarrow$ Dw (down + pen)
2	First pt. of positive label	Up $\rightarrow$ Up (null), Dw $\rightarrow$ Dw (null), Dw $\rightarrow$ nCU (up + pen), Up $\rightarrow$ nCD (down + pen)
3	Interior of positive label	Up $\rightarrow$ Up (null), Dw $\rightarrow$ Dw (null), nCU $\rightarrow$ nCU (null), nCD $\rightarrow$ nCD (null), Dw $\rightarrow$ nCU (up + pen), Up $\rightarrow$ nCD (down + pen)
4	End of positive label	nCU $\rightarrow$ Up (null), nCD $\rightarrow$ Dw (null), Dw $\rightarrow$ Up (up + pen), Up $\rightarrow$ Dw (down + pen)
5	Single-point positive label	Dw $\rightarrow$ Up (up + pen), Up $\rightarrow$ Dw (down + pen)
6	In negative label	Up $\rightarrow$ Up (null), Dw $\rightarrow$ Dw (null)

Table 3: *Up-down with labels: 4 states $\times$ 6 rules. Rules 2 and 5 deliberately exclude noChange self-loops to avoid infinity arithmetic.*

Key design insight: in the standard up-down model, every changepoint is a direction change (Up $\rightarrow$ Dw or Dw $\rightarrow$ Up). So requiring *at least one changepoint* in a positive label is equivalent to requiring at least one direction change. Rule 4 exploits this: the only exits from the Up/Dw states are direction changes, which is not an over-constraint but a natural consequence of the up-down structure.

Rules 2 and 5 exclude noChange self-loops that would operate on infinite-cost `ListPieces` at unreachable states. Rule 3 includes these self-loops because by the interior of a positive label, the noChange states may have been reached via a change transition. The 6-rule decomposition is my current best design; I’m curious whether the mentors see a simpler formulation.

Validation: No oracle exists. Instead:

- Construct synthetic labeled genomic data with known peaks.
- Run `gfpop` with the 6-rule graph.
- Verify structural invariants: exactly one changepoint in each positive label region, zero changepoints in each negative label region.
- Compare unlabeled-region behavior against the standard `gfpop` up-down model (should match when all rules are set to the unlabeled rule).
- Hand-trace the 6-rule model on a small (5–10 point) dataset as part of the design doc, verifying cost optimality before writing code.

2.8.3 Capstone Deliverable: Vignette and Visualization

A formal, user-facing R vignette that:

1. Introduces the problem: “You have labeled data. Here’s how to use it with `gfpop`.”
2. Walks through a genomic peak detection example with labels.
3. Shows the `plotModel()` visualization of the rule-partitioned graph.
4. Compares `gfpop` output with LOPART output on the same data.
5. Includes a **filmstrip plot**: a sequence of panels showing the active subgraph at representative time points (one per rule ID).

The filmstrip visualization will be a formal deliverable. I found that the hardest part of learning gfpop was understanding which edges fire when, and this is the diagram I wished existed.

2.8.4 Testing and Benchmarking Strategy

Five layers of testing:

- **Oracle testing:** `gfpop(rule = LOPART_graph)` vs `LOPART::LOPART()`: exact match on segments, means, and total cost across multiple test datasets and penalties.
- **Regression testing:** `gfpop()` without `rule` produces bit-identical results to current gfpop on the existing test suite.
- **Structural invariant testing:** For the up-down+labels model, verify exactly one change-point per positive label region and zero per negative label region on synthetic data.
- **Fuzz testing:** Random data + random label positions, check that `gfpop(rule=...)` never crashes, produces NaN, or violates label constraints. 200+ trials.
- **Benchmarking:** `microbenchmark` comparing `gfpop(rule = LOPART)` vs standalone `LOPART::LOPART()` at $N = 10^3, 10^4, 10^5, 10^6$. Report median wall-clock time. Standalone LOPART has a tightly optimized loop for its specific model, so gfpop's general graph machinery may be slower by a constant factor. The goal is to quantify that overhead and confirm it stays within a reasonable range (e.g., $< 5\times$).

2.9 Success Criteria and Non-Goals

2.9.1 Core Goals (must deliver)

- `Edge(..., rule=)` parameter added to R API, backward compatible
- `gfpop(..., rule=)` parameter added, backward compatible
- C++ edge filtering by rule in `LP_edges_operators`
- Infinity handling in `LP_ts_Minimization` and `LP_edges_addPointAndPenalty`
- LOPART model (4 rules, 2 states) working in gfpop, oracle-validated against `LOPART::LOPART()`
- Regression tests: gfpop without `rule` produces identical results to current gfpop
- R CMD check passes with no new warnings

2.9.2 Stretch Goals (if time permits)

- Up-down with labels (6 rules, 4 states) working in gfpop
- Performance benchmarks: `gfpop(rule=...)` vs standalone LOPART, $N = 10^3$ to 10^6
- Capstone vignette with filmstrip visualization
- Design a FLOPART-equivalent model graph (Poisson loss, up-down constraints, labels). This is a model design task, not a C++ change.
- RLE compression of the rule vector for long constant-label runs

2.9.3 Non-Goals

- **Not modifying the PDPA algorithm.** Only the FPOP (penalty-based) solver.
- **Not redesigning gfpop's cost function dispatch.** gfpop uses global function pointers rather than a class hierarchy. I will work within this pattern, not refactor it.

2.10 Timeline

Phase	Week	Details	Hours
Pre-acceptance	Now – Apr 30	Communicate with mentors, study gfpop and FLOPART C++ source (especially infinity handling at label boundaries), prototype infinity sentinel	N/A
Community Bonding	May 1–7	Fork gfpop, set up dev environment, configure CI	N/A
	May 8–14	Write design doc with mentors’ input, finalize API	N/A
	May 15–24	Implement <code>Edge(rule=)</code> and <code>graph(rule=)</code> on R side	N/A
Part 1	May 25–31	Add <code>ruleID</code> to <code>Edge</code> , rule vector to <code>Omega</code> , pre-computed lookup	30
	Jun 1–7	Edge filtering in <code>LP_edges_operators</code> , infinity sentinel	30
	Jun 8–14	Unit tests: infinity through <code>operatorUp</code> , <code>operatorDw</code> , min-envelope	25
	Jun 15–21	Build LOPART graph in R, extend backtracking for rule-aware <code>findBeta</code> , first end-to-end test	30
	Jun 22–28	Oracle tests against <code>LOPART::LOPART()</code> , debug	35
	Jun 29–Jul 5	Hardening: edge cases, backward compat. Midterm report	30
	Jul 6–10	Midterm: LOPART model working, oracle tests passing	
Part 2	Jul 6–12	Design 6-rule up-down graph, enumerate edges per rule	25
	Jul 13–19	Implement up-down+labels, validate on synthetic data	30
	Jul 20–26	Test label constraint satisfaction, 4-state edge cases	25
	Jul 27–Aug 2	Performance benchmarks, begin capstone vignette	25
	Aug 3–9	Extend <code>plotModel</code> for 6-rule graph, filmstrip visualization	20
	Aug 10–16	Polish: man pages, comprehensive test suite (fuzz, large N)	25
	Aug 17–24	Final code review with mentors, address feedback, submit PR	15
	Aug 17–31	Final evaluation	
Total			~345

Contingency: The vignette and visualization weeks (Jul 27 – Aug 9) are buffer. If the core implementation takes longer, the vignette scope shrinks but the algorithm ships. If ahead of schedule, extend to Poisson loss with labels.

If behind schedule: The 6-rule up-down model becomes a stretch goal. LOPART (4 rules) with full test coverage and benchmarks is the minimum viable deliverable.

Balancing coursework: I will be taking summer courses alongside GSoC. The community bond-

ing period is deliberately front-loaded (fork, CI, design doc, R-side API) so that Part 1 coding begins with infrastructure already in place. The heavier coding weeks (30 hrs) are in June before course deadlines peak; the lighter weeks (15–20 hrs) are in August when I expect more exams.

2.11 Management of Coding Project

- **Commits:** Daily pushes to a feature branch. PRs for each milestone:
 - (1) R API + graphReorder, (2) C++ infrastructure + infinity,
 - (2) LOPART model, (4) up-down model, (5) vignette.
- **CI:** R CMD check on GitHub Actions for every PR. The full testthat suite must pass before merge.
- **Communication:** Weekly status update to mentors via email or GitHub issue thread. If no commits for 3+ days, proactively flag blockers.
- **Code review:** Request mentor review at each milestone before proceeding to the next.
- **Red flag protocol:** If commits slow to fewer than 2 per week, or tests fail without quick fixes, escalate to mentors immediately.

3 About Me

3.1 Bio

Hi! I'm William. I study Mathematics at the University of Waterloo, Canada (Statistics major, Computing minor). At the core, I'm a stats student, so I use R quite often. I also work in actuarial science, currently interning at Intact (previously at Pacific Life Re). But the part I keep coming back to is the modeling.

At Intact, I got curious about usage-based insurance: telematics pricing, where you're trying to figure out how someone's driving behavior maps to their risk. That's a changepoint problem. You're looking for the moment a driver's pattern shifts, a trend break in claim frequency, an inflection point in how losses develop. The more I explored the newer corners of actuarial work, the more I kept finding problems that looked like constrained changepoint detection. That's how I found gfpop.

This project makes sense to me. The extension is small and surgical, but it lets gfpop express every labeled changepoint model. That's an elegant payoff. On a personal level, it also connects to the applied problems I care about. So ya, I want to build this because it's cool, and because people like me would actually use it.

3.2 Programming Languages and Tools

I picked up R through my statistics courses and C++ through my computing minor. R is what I reach for first for anything statistics-related, and C++ is what the qualification tests pushed me deeper into. I wrote roughly 1,200 lines across three solvers (piecewise cost operators, linked-list manipulation, memory management), I'm still learning C++, but the tests taught me to read an unfamiliar numerical codebase, trace the logic, and get working changes into it.

On the R side, I'm comfortable with ggplot2, testthat, Rcpp, and the CRAN development pipeline (NAMESPACE, .C() interface, C registration, R CMD check).

Python, SQL, and Next.js come from the work side: internships at Intact and Pacific Life Re, plus hackathons and a research project. And LaTeX is what this proposal is written in.

3.3 Skills I Picked Up

The qualification tests (see Section~1.6) taught me a lot. Going in, I knew some R and had a bit of C++ from coursework. Coming out, I had a better sense of what it actually takes to work on numerical software.

The biggest lesson was learning to read someone else's C++ codebase. PeakSegOptimal is roughly 3,000 lines with no tutorial. I spent a lot of time tracing the DP loop and mapping out the flow before I changed anything. Building a shared LossPiece base class for NormalLossPiece and PoissonLossPieceLog gave me a small taste of what gfpop's multi-loss architecture looks like from the inside. That's the same skill I'll need for gfpop's Omega, ListPiece, and Graph classes.

I also learned that "it runs" is not the same as "it's correct." Oracle tests (against Segmentor3IsBack, isoreg(), fpop::Fpop()) caught subtle off-by-one and boundary errors that my unit tests missed entirely. Fuzz testing (200-trial randomized inputs) surfaced edge cases I hadn't thought to check.

I plan to use both heavily for this project. And integrating everything into a CRAN-ready package (C registration, `.C()` interface, `testthat`, `R CMD check`) taught me the plumbing that papers never talk about.

3.4 What I Expect from GSoC

Mostly, I want to learn what it takes to ship real code into a package that people actually use. I've written solvers on my own, but contributing to someone else's codebase is different. Navigating existing design decisions, writing code that a maintainer would actually want to merge. The PR review process gave me a taste of that, and it's a skill I haven't practiced enough.

I'm also curious about the research-to-software pipeline. The ideas behind this extension already exist in papers (LOPART, FLOPART, `gfpop`), but turning them into a clean, tested, backward-compatible implementation involves decisions that papers never talk about. I want to get better at making those calls.

Beyond this summer, I'd like to stay involved with `gfpop`. The usage-based side of actuarial science, where I want to work, treats changepoint detection as a recurring part of the modeling workflow, not a one-off analysis. Staying close to `gfpop` means I'd be both building the tool and using it. That feels like the right way to contribute.

That, and I really want to debug infinity propagation through `operatorUp`. It's quite cool.

4 References

- Jackson, B., Scargle, J.D., Barnes, D., Arabhi, S., Alt, A., Gioumoussis, P., Gwin, E., Sangtrakulcharoen, P., Tan, L., and Tsai, T.T. (2005). “An Algorithm for Optimal Partitioning of Data on an Interval.” *IEEE Signal Processing Letters*, 12(2), 105–108.
- Kaufman, J.M., Stenberg, A.J., and Hocking, T.D. (2024). “Functional Labeled Optimal Partitioning.” *Journal of Computational and Graphical Statistics*. <https://doi.org/10.1080/10618600.2023.2293216>
- Maidstone, R., Hocking, T.D., Rigai, G., and Fearnhead, P. (2017). “On Optimal Multiple Changepoint Algorithms for Large Data.” *Statistics and Computing*, 27(2), 519–533. <https://doi.org/10.1007/s11222-016-9636-3>
- Runge, V., Hocking, T.D., Romano, G., Afghah, F., Fearnhead, P., and Rigai, G. (2020). “gfpop: an R Package for Univariate Graph-Constrained Change-point Detection.” <https://arxiv.org/abs/2002.03646>
- Srivastava, A. and Hocking, T.D. (2023). “LOPART: Labeled Optimal Partitioning.” *Computational Statistics*, 38, 461–480. <https://doi.org/10.1007/s00180-022-01238-z>
- gfpop. <https://github.com/vrunge/gfpop>
- PeakSegOptimal. <https://github.com/tdhock/PeakSegOptimal>